



Tudor Cebere

PhD Student, Inria

tudor.cebere@inria.fr

tudorcebere.github.io

Differential Privacy & Auditing for Machine Learning

Final-year PhD student at **Inria**, advised by **Aurélien Bellet**, and recipient of the **Google PhD Fellowship 2025** in **Privacy, Safety, and Security**.

Expected to graduate in **December 2026**. In Fall 2026, I will join **Google** as a research intern, working on **private analytics for large models**.

On the job market starting early 2027.

Research Interests

- **Privacy auditing.** I make privacy audits realistic and actionable: fully observational auditing [1], practical DP-SGD threat models [2], and tools for real DP libraries [3]. My longer-term goal is to facilitate measuring privacy lower bounds across algorithms, settings, making risk measurable and communicable, particularly for large models.
- **Private learning in practice.** I study DP learning under realistic deployment constraints, including varied threat models [2] and constrained learning settings [4].
- **Current interest.** I am aiming a) continue to push the limits of what we can do with privacy auditing (particularly for machine learning) and b) seeking for alternatives for differentially private machine learning beyond DP-SGD.

Selected work

[1] **Privacy Auditing with Zero Training Run.** Cebere*, Even*, Bleistein, Bellet, Preprint, 2026.

[3] **Grey-Box Auditing of DP Libraries.** Cebere, Erb, Desfontaines, Bellet, Fitzsimons, PETS 2026.

[2] **Tighter Privacy Auditing of DP-SGD.** Cebere, Bellet, Papernot, ICLR 2025.

[4] **Private Rate-Constrained Optimization.** Yaghini*, Cebere*, Menart, Bellet, Papernot, ICLR 2026.